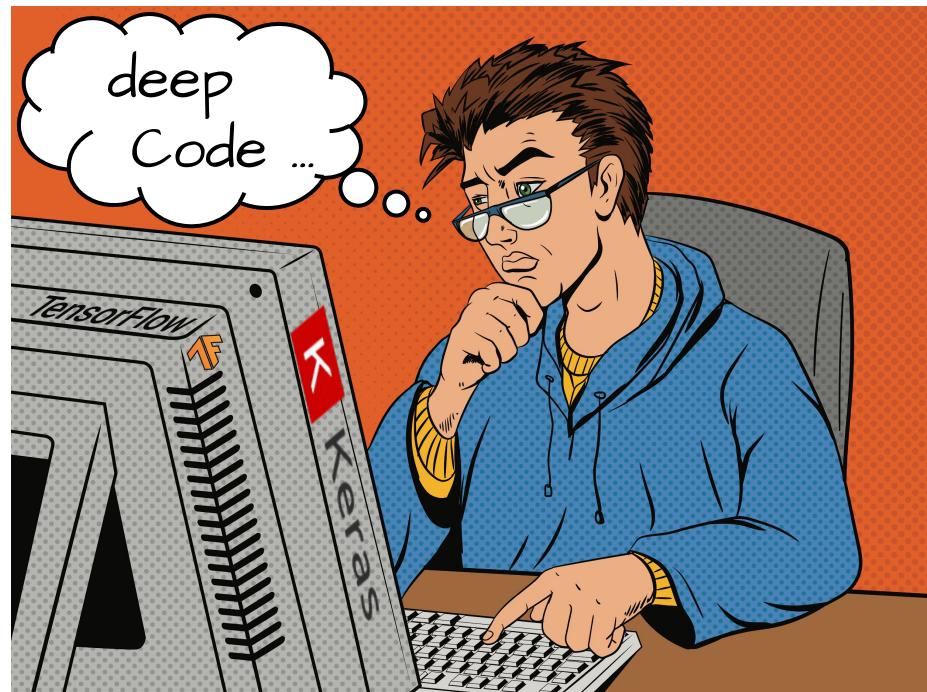


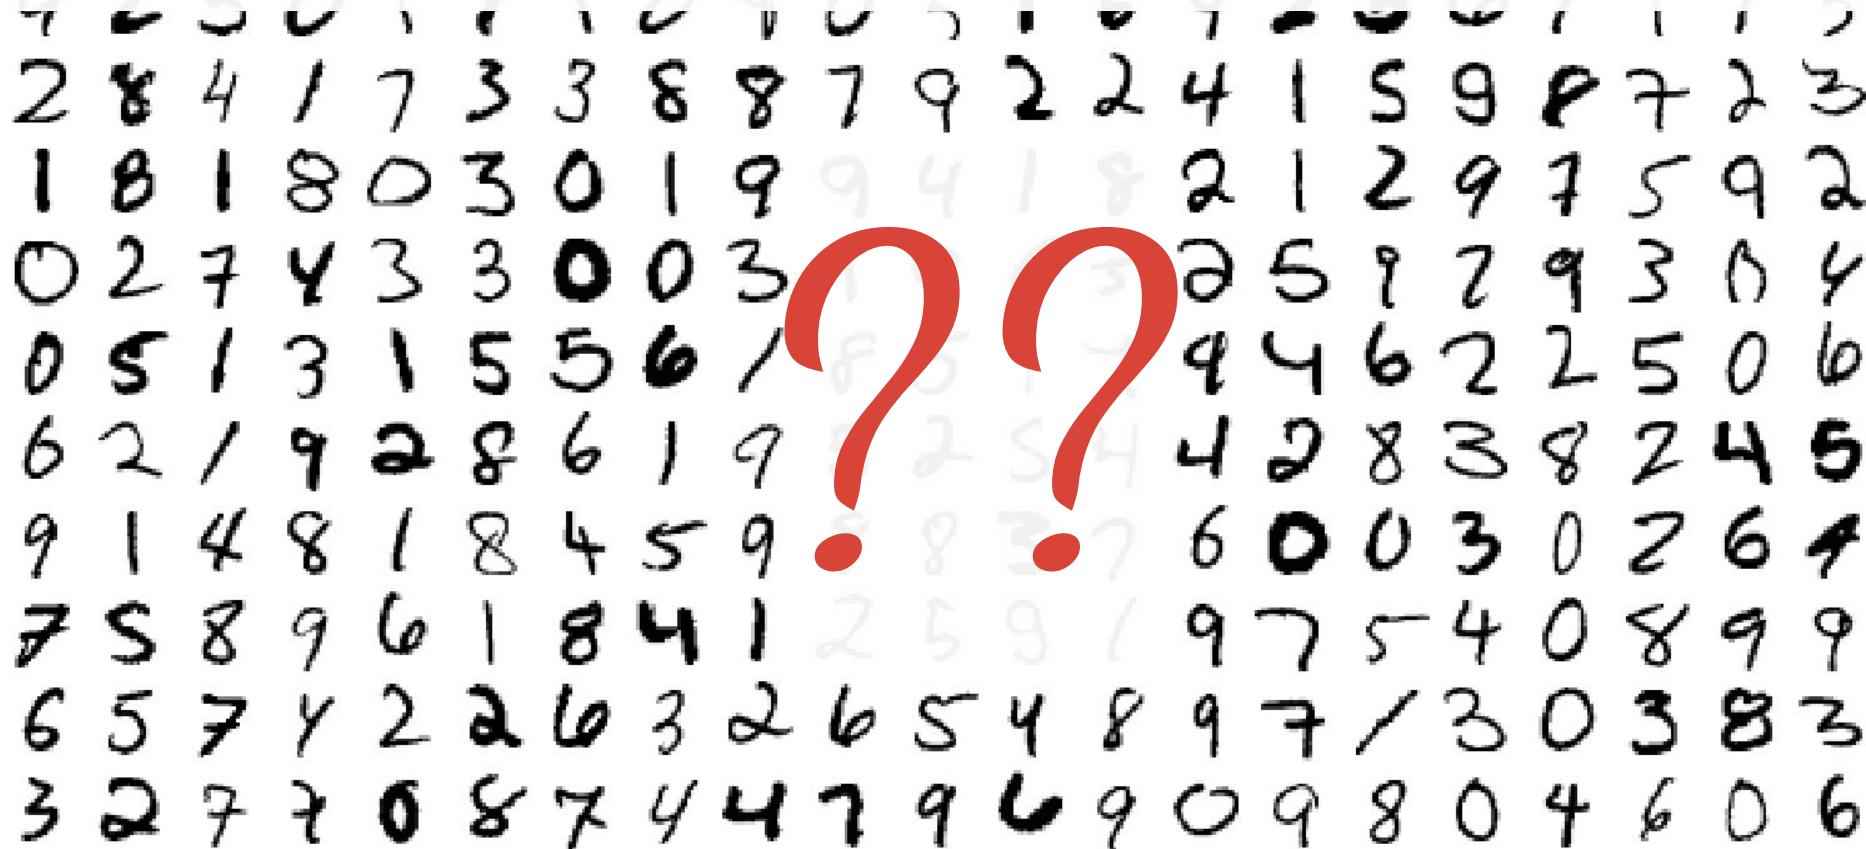
>Optimising Training of Neural Networks_



Recap till now

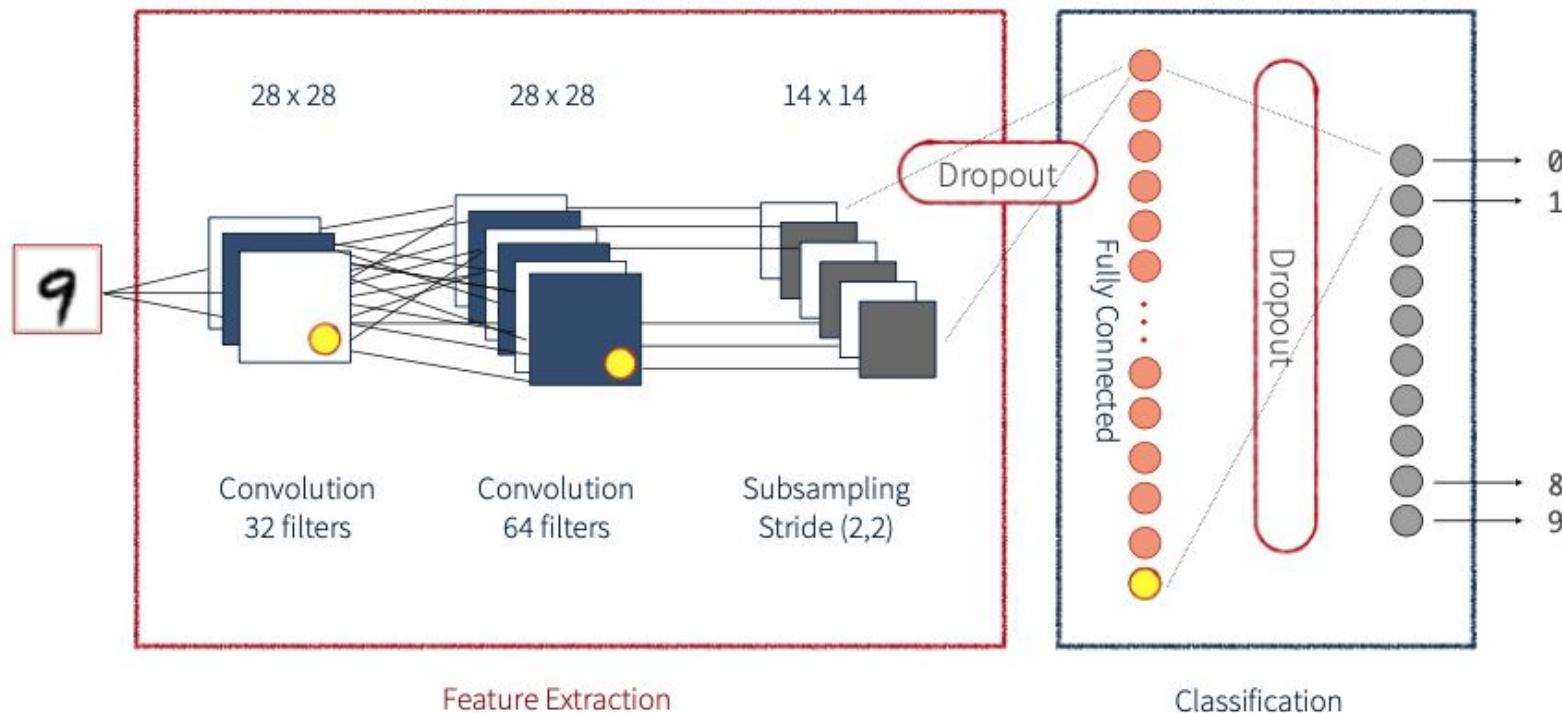
- Introduction to Neural Networks & its working
- Tensorflow framework
- Introduction to Neural Networks for Regression Problems
- Introduction to Neural Networks for Classification Problems
- Introduction to Hyperparameters (epochs, learning rate, batch size)
- Basics of optimization

Hello World: handwritten digits classification - MNIST

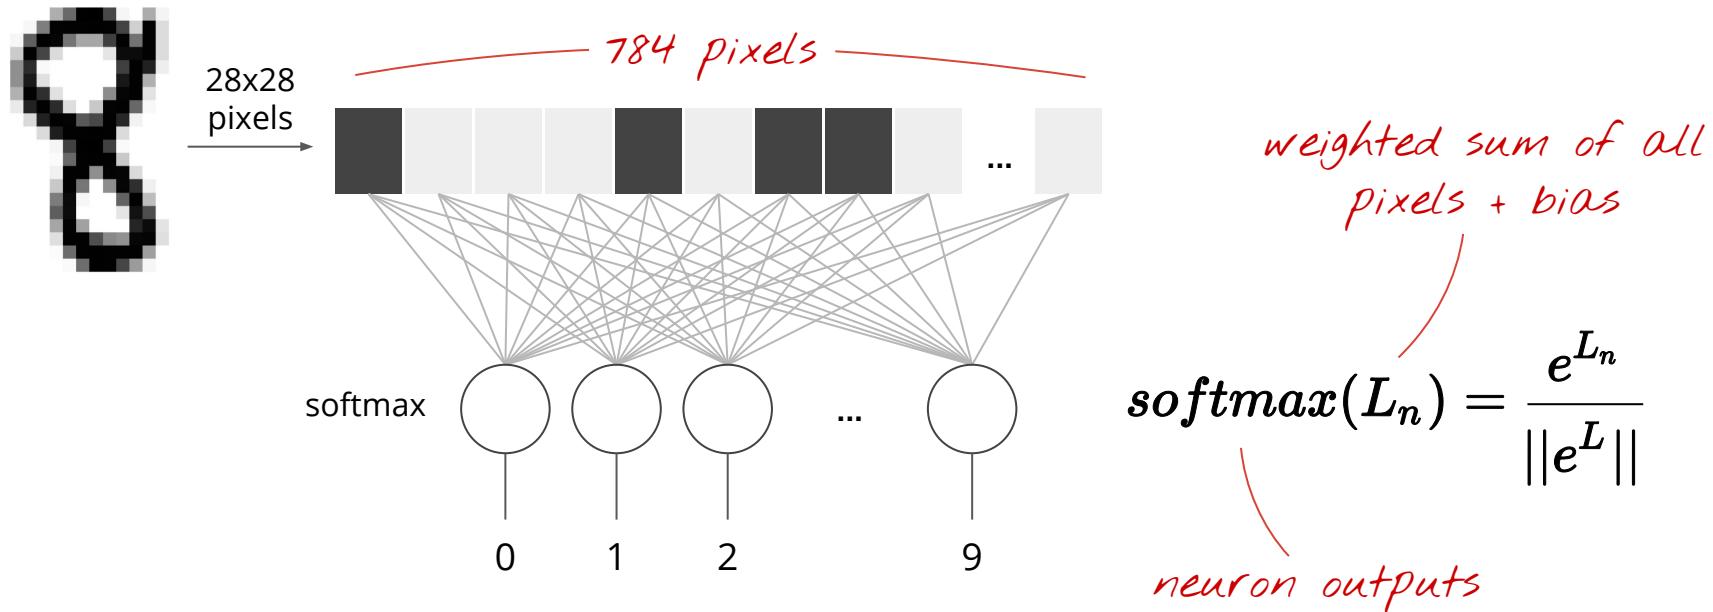


MNIST = Mixed National Institute of Standards and Technology - Download the dataset at <http://yann.lecun.com/exdb/mnist/>

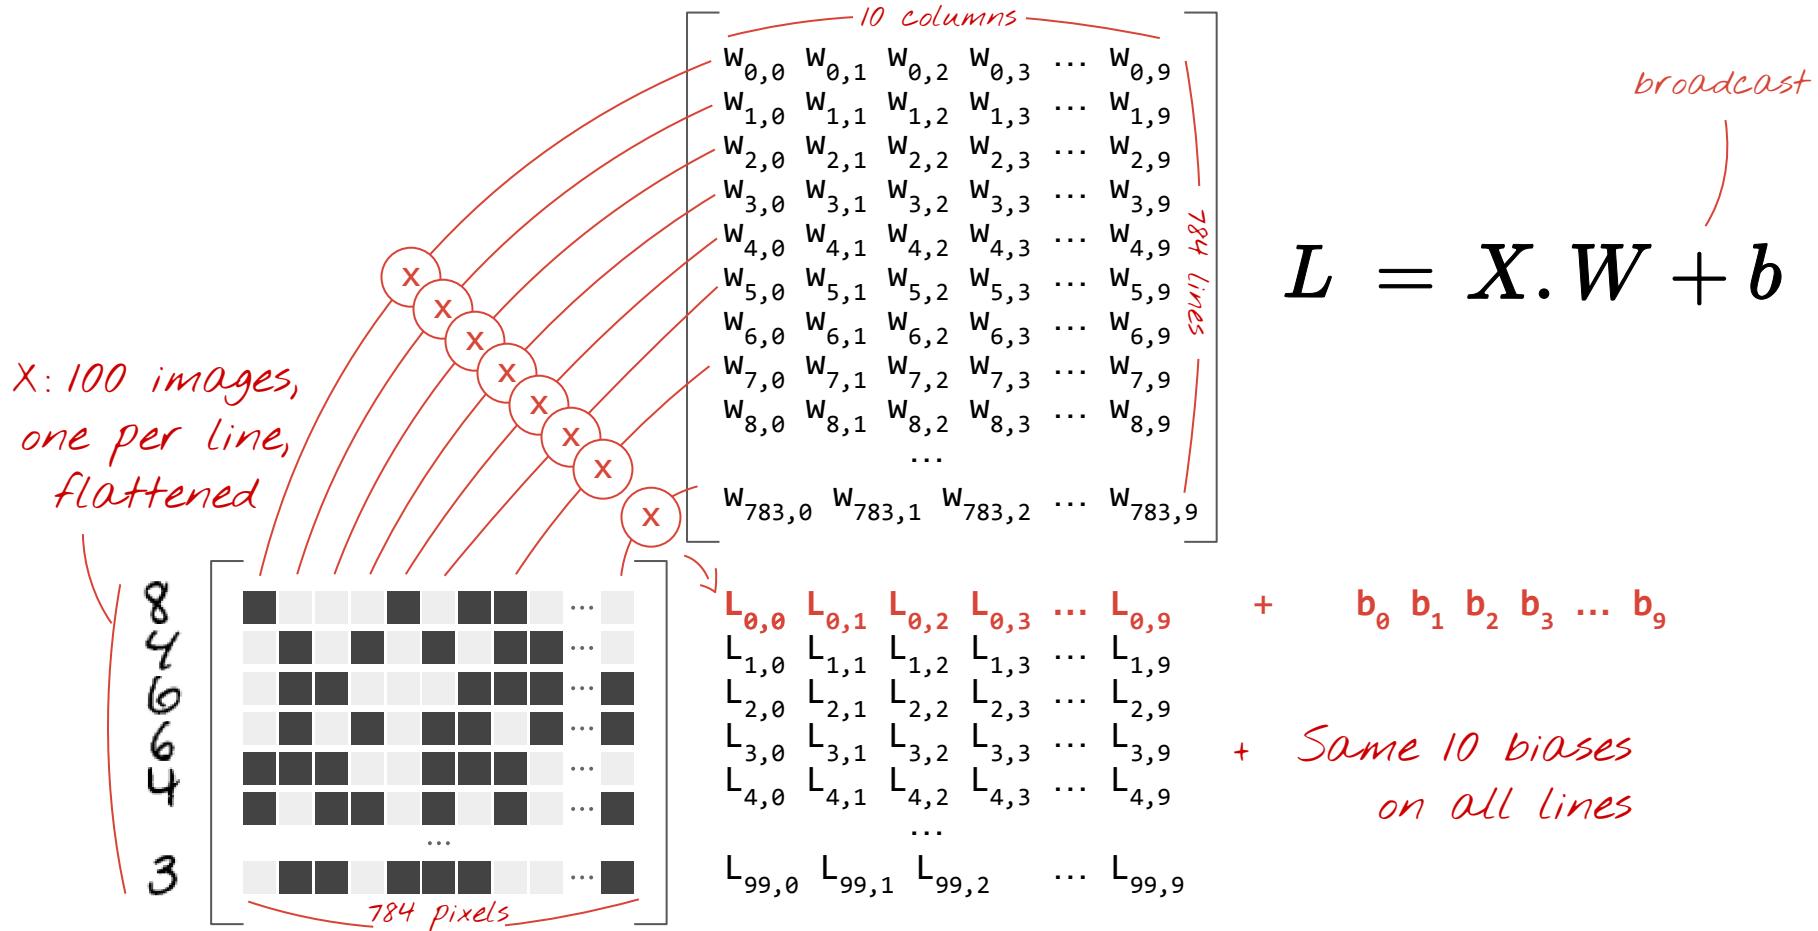
Convolutional Neural Networks



Very simple model: softmax classification



In matrix notation, 100 images at a time



Softmax, on a batch of images

$$Y = \text{softmax}(X \cdot W + b)$$

Predictions
 $Y[100, 10]$

Images
 $X[100, 784]$

Weights
 $W[784, 10]$

Biases
 $b[10]$

applied line by line

matrix multiply

broadcast on all lines

tensor shapes in []

Now in TensorFlow (Python)

tensor shapes: $X[100, 784]$ $W[748,10]$ $b[10]$

$Y = \text{tf.nn.softmax}(\text{tf.matmul}(X, W) + b)$

matrix multiply *broadcast
on all lines*

Success ?

0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	1	0	0	0

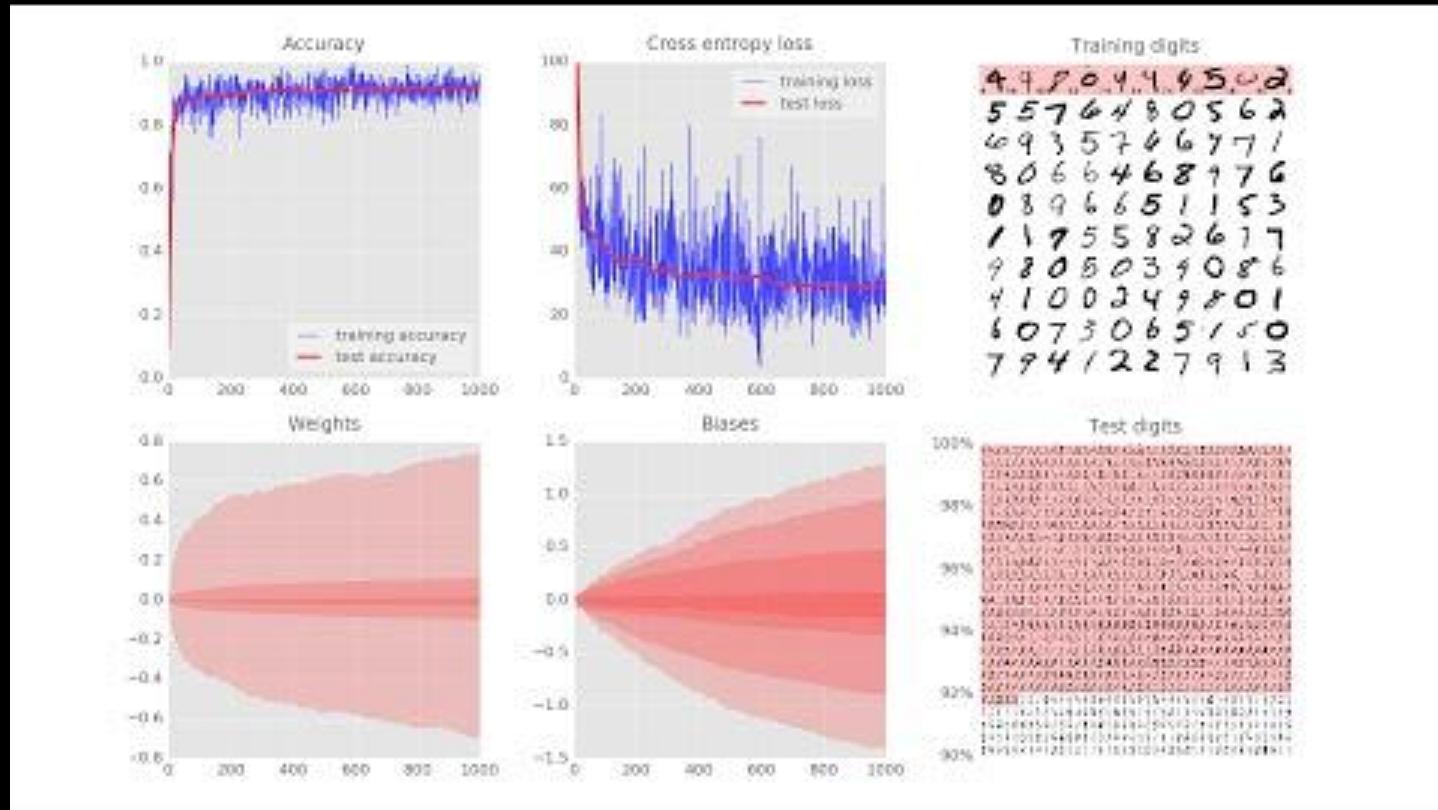
actual probabilities, "one-hot" encoded

Cross entropy: $-\sum Y'_i \cdot \log(Y_i)$

)
this is a "6"
computed probabilities

.01	.03	.00	.04	.03	.05	0.8	.02	.01	.01
0	1	2	3	4	5	6	7	8	9

Demo



TensorFlow - initialisation

```
import tensorflow as tf
```

this will become the batch size, 100

/

```
X = tf.placeholder(tf.float32, [None, 28, 28, 1])
```

```
W = tf.Variable(tf.zeros([784, 10]))
```

```
b = tf.Variable(tf.zeros([10]))
```

\

28 x 28 grayscale images

```
init = tf.initialize_all_variables()
```

Training = computing variables W and b

TensorFlow - success metrics

```
# model  
Y = tf.nn.softmax(tf.matmul(tf.reshape(X, [-1, 784]), W) + b)  
# placeholder for correct answers  
Y_ = tf.placeholder(tf.float32, [None, 10])  
# loss function  
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))  
# % of correct answers found in batch  
is_correct = tf.equal(tf.argmax(Y, 1), tf.argmax(Y_, 1))  
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

flattening images

/

"one-hot" encoded

"one-hot" decoding

TensorFlow - training

```
optimizer = tf.train.GradientDescentOptimizer(0.003)  
train_step = optimizer.minimize(cross_entropy)
```

learning rate



loss function



TensorFlow - run !

```
sess = tf.Session()  
sess.run(init)  
  
for i in range(1000):  
    # Load batch of images and correct answers  
    batch_X, batch_Y = mnist.train.next_batch(100)  
    train_data={X: batch_X, Y_: batch_Y}  
  
    # train  
    sess.run(train_step, feed_dict=train_data)
```

running a Tensorflow computation, feeding placeholders

```
# success ?  
a,c = sess.run([accuracy, cross_entropy], feed_dict=train_data)  
  
# success on test data ?  
test_data={X: mnist.test.images, Y_: mnist.test.labels}  
a,c = sess.run([accuracy, cross_entropy], feed=test_data)
```

Tip:

do this
every 100
iterations

TensorFlow - full python code

```
import tensorflow as tf

X = tf.placeholder(tf.float32, [None, 28, 28, 1])
W = tf.Variable(tf.zeros([784, 10]))
b = tf.Variable(tf.zeros([10]))
init = tf.initialize_all_variables()

# model
Y = tf.nn.softmax(tf.matmul(tf.reshape(X, [-1, 784]), W) + b)

# placeholder for correct answers
Y_ = tf.placeholder(tf.float32, [None, 10])

# loss function
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))

# % of correct answers found in batch
is_correct = tf.equal(tf.argmax(Y, 1), tf.argmax(Y_, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

initialisation

model

success metrics

training step

```
optimizer = tf.train.GradientDescentOptimizer(0.003)
train_step = optimizer.minimize(cross_entropy)

sess = tf.Session()
sess.run(init)

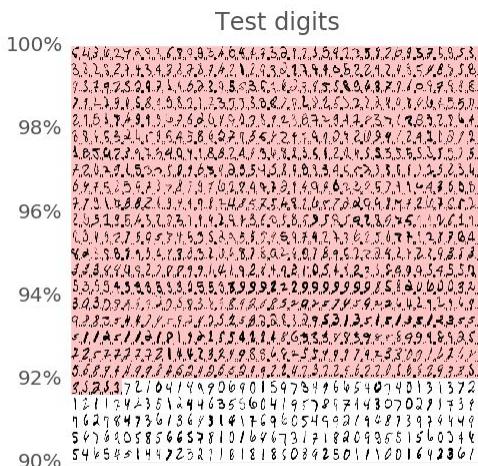
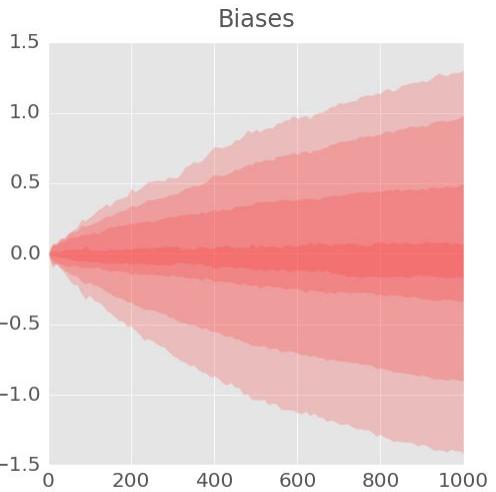
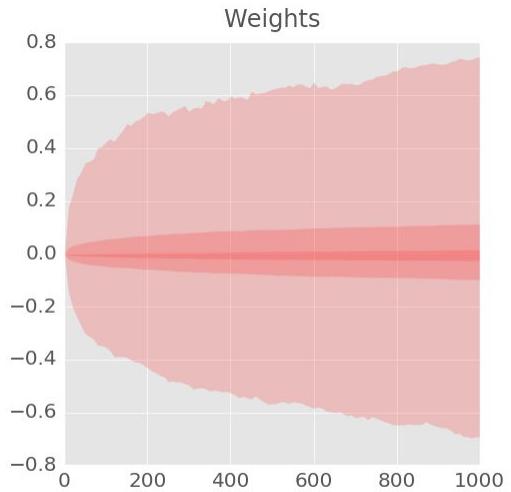
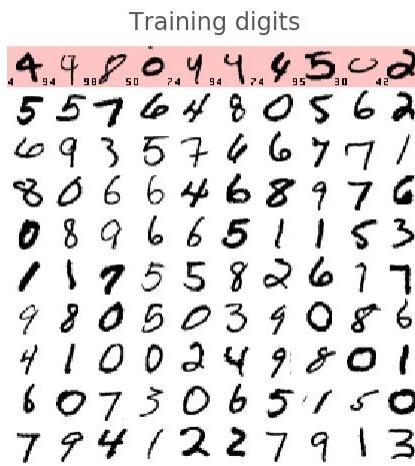
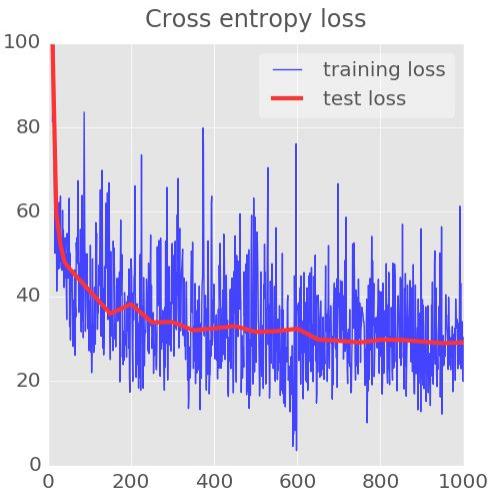
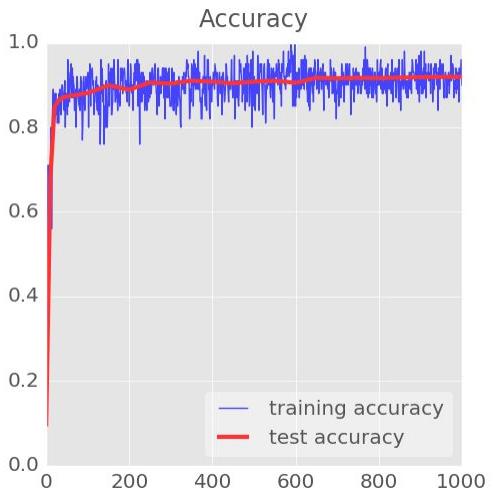
for i in range(10000):
    # load batch of images and correct answers
    batch_X, batch_Y = mnist.train.next_batch(100)
    train_data = {X: batch_X, Y_: batch_Y}

    # train
    sess.run(train_step, feed_dict=train_data)

    # success ? add code to print it
    a, c = sess.run([accuracy, cross_entropy], feed=train_data)

    # success on test data ?
    test_data = {X: mnist.test.images, Y_: mnist.test.labels}
    a, c = sess.run([accuracy, cross_entropy], feed=test_data)
```

Run



92%

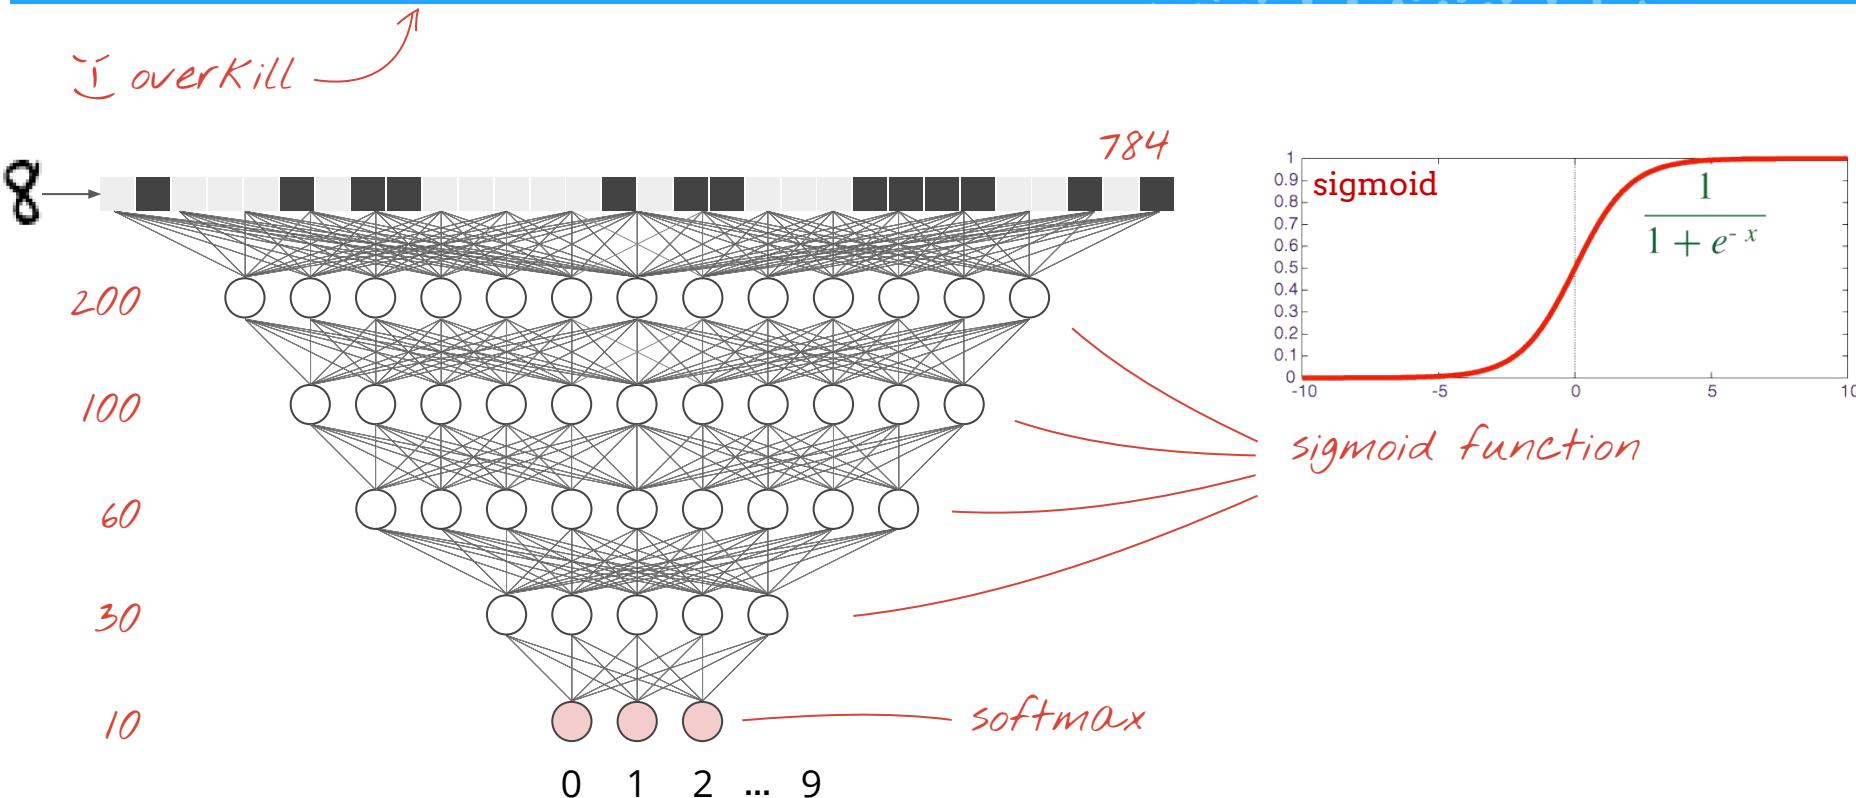
Softmax
Cross-entropy
Mini-batch





Go deep !

Let's try 5 fully-connected layers !



TensorFlow - initialisation

K = 200
L = 100
M = 60
N = 30

*weights initialised
with random values*

```
W1 = tf.Variable(tf.truncated_normal([28*28, K] ,stddev=0.1))  
B1 = tf.Variable(tf.zeros([K]))
```

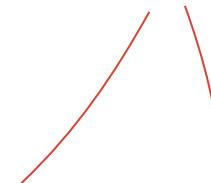
```
W2 = tf.Variable(tf.truncated_normal([K, L], stddev=0.1))  
B2 = tf.Variable(tf.zeros([L]))
```

```
W3 = tf.Variable(tf.truncated_normal([L, M], stddev=0.1))  
B3 = tf.Variable(tf.zeros([M]))  
W4 = tf.Variable(tf.truncated_normal([M, N], stddev=0.1))  
B4 = tf.Variable(tf.zeros([N]))  
W5 = tf.Variable(tf.truncated_normal([N, 10], stddev=0.1))  
B5 = tf.Variable(tf.zeros([10]))
```

TensorFlow - the model

```
X = tf.reshape(X, [-1, 28*28])
```

weights and biases



```
Y1 = tf.nn.sigmoid(tf.matmul(X, W1) + B1)
```

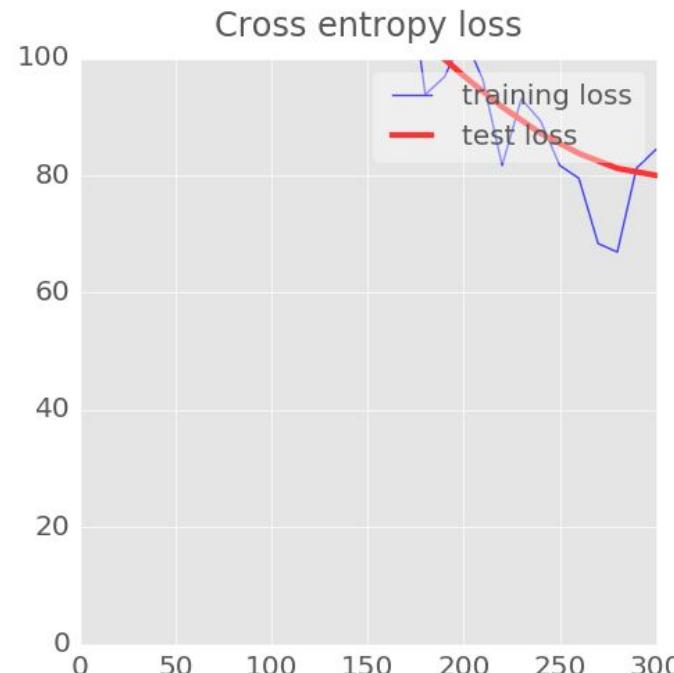
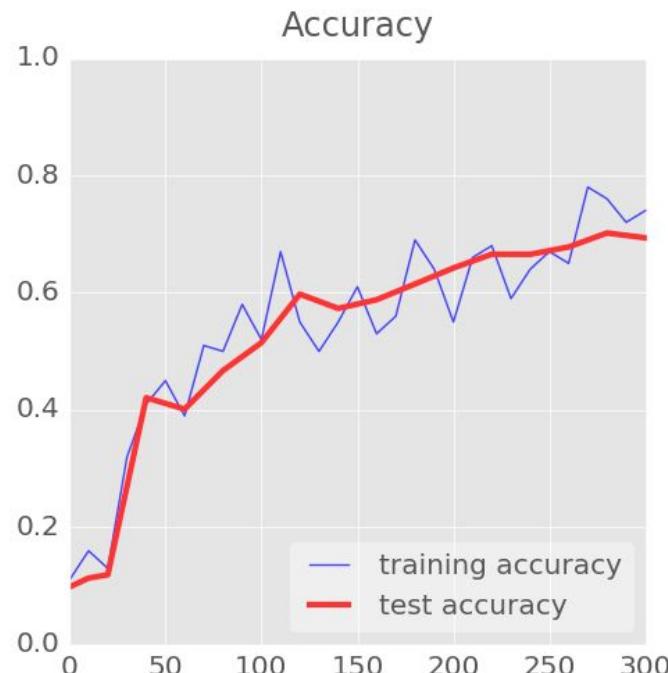
```
Y2 = tf.nn.sigmoid(tf.matmul(Y1, W2) + B2)
```

```
Y3 = tf.nn.sigmoid(tf.matmul(Y2, W3) + B3)
```

```
Y4 = tf.nn.sigmoid(tf.matmul(Y3, W4) + B4)
```

```
Y = tf.nn.softmax(tf.matmul(Y4, W5) + B5)
```

slow start?



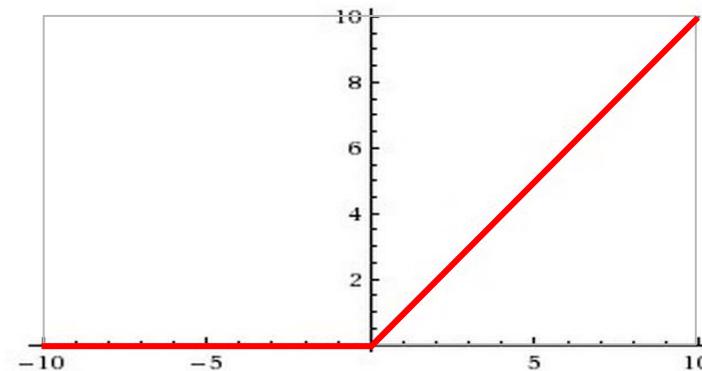
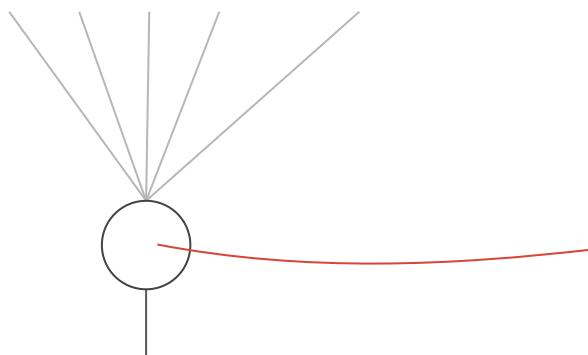


Relu !

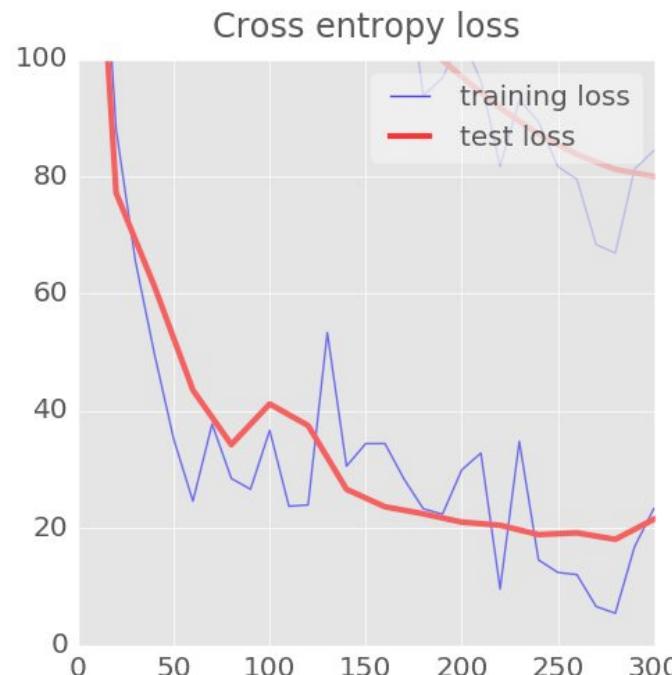
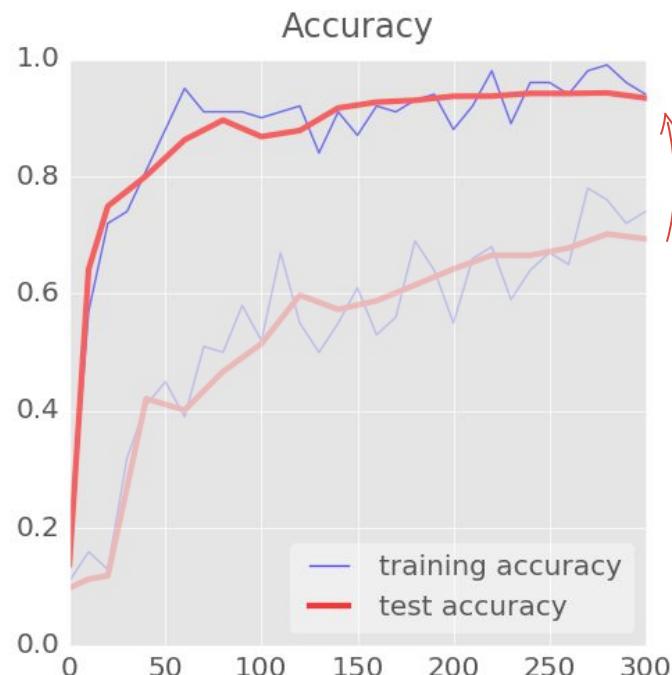


RELU

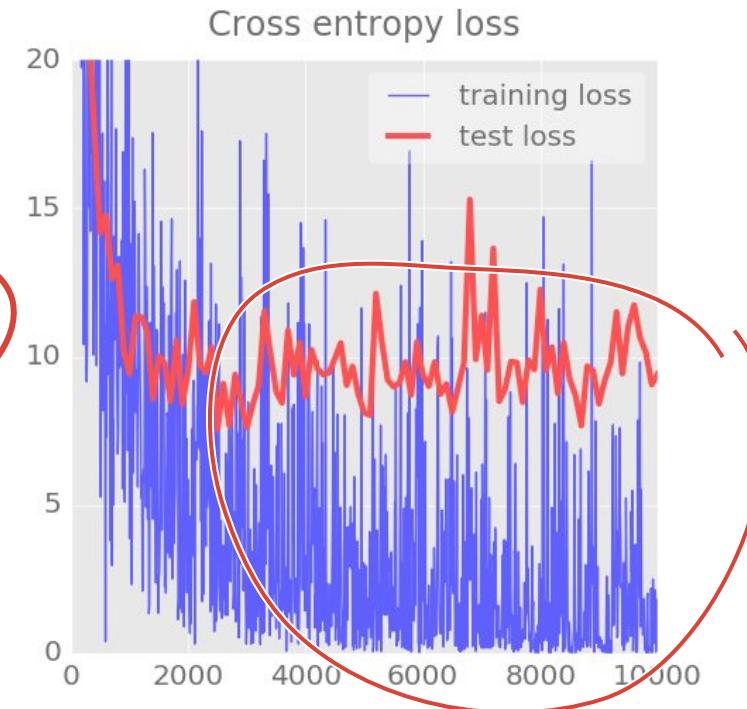
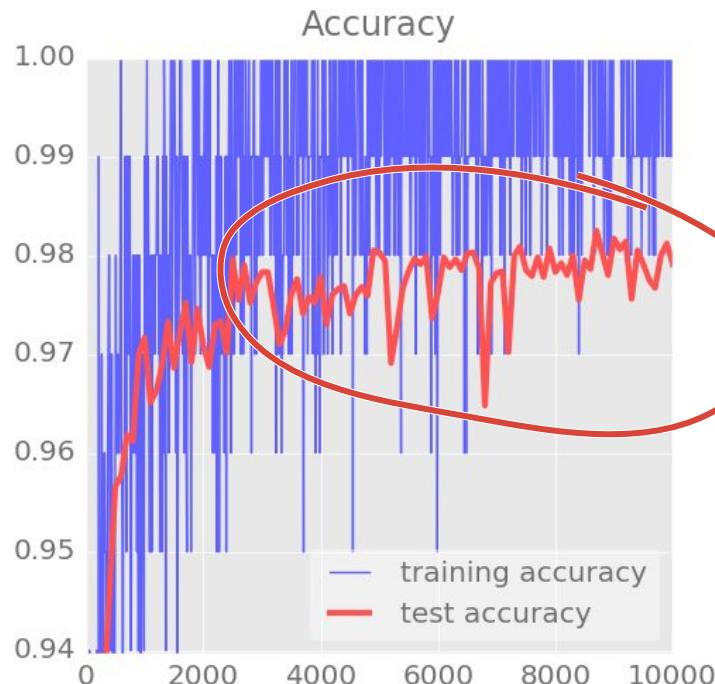
RELU = Rectified Linear Unit


$$Y = \text{tf.nn.relu}(\text{tf.matmul}(X, W) + b)$$

RELU



Demo - noisy accuracy curve ?



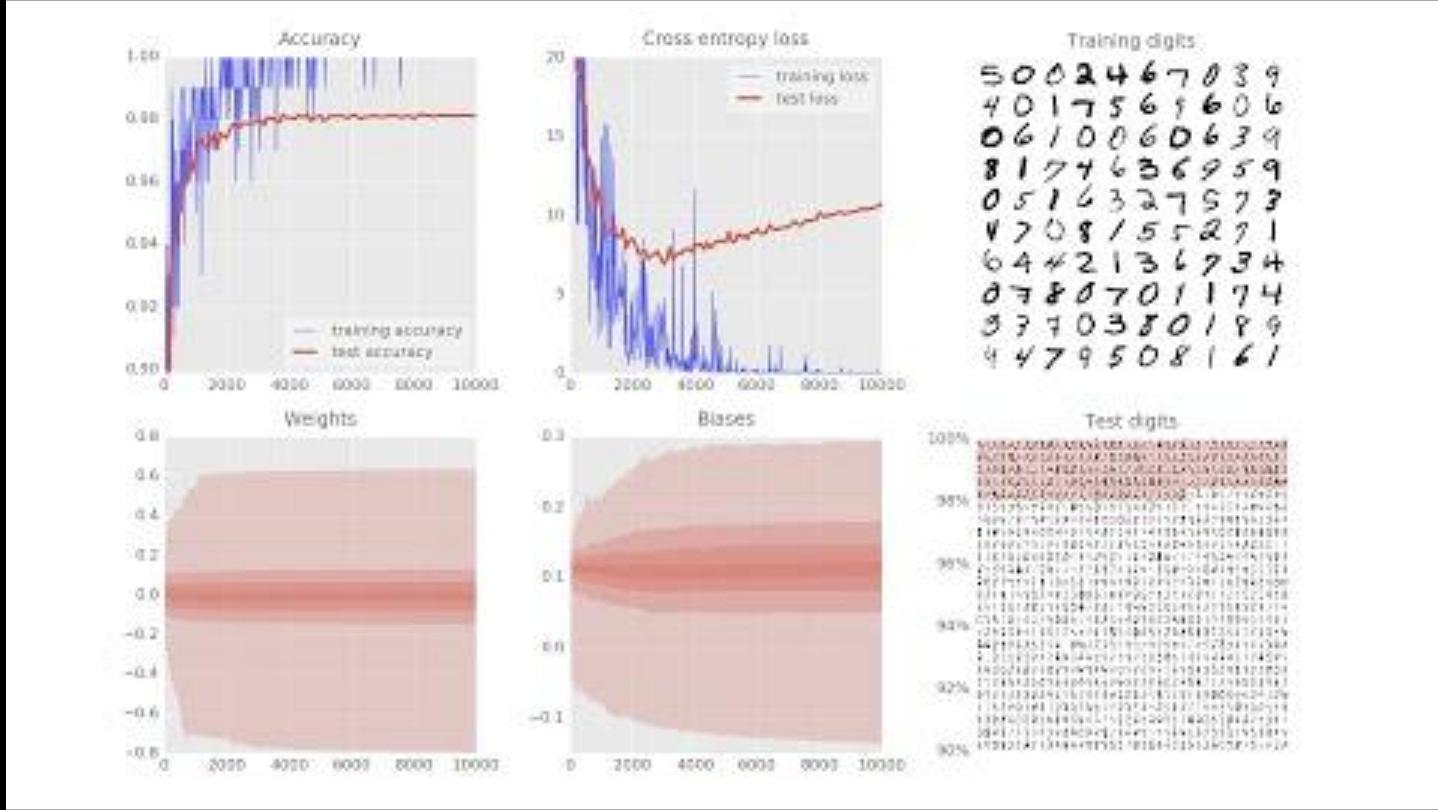
yuck!

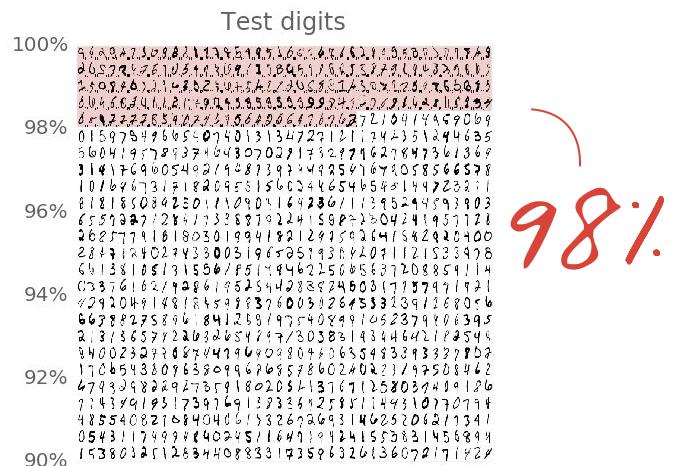
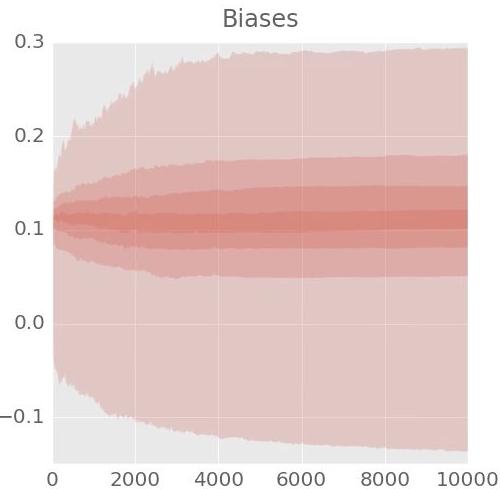
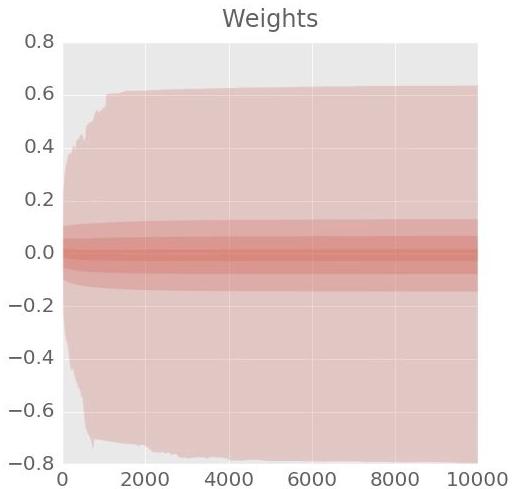
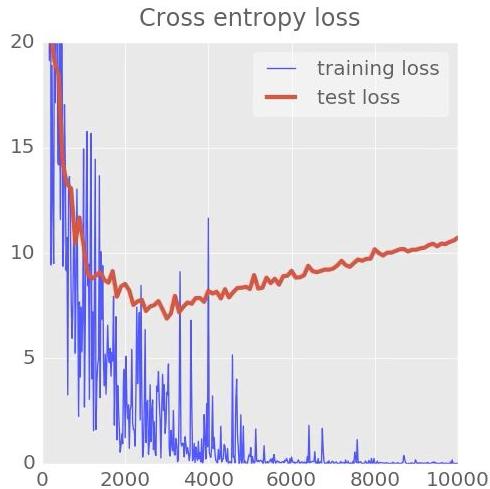
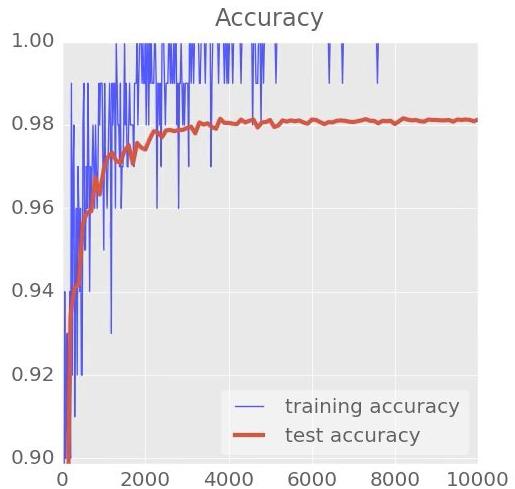
Slow down...

Learning
rate decay

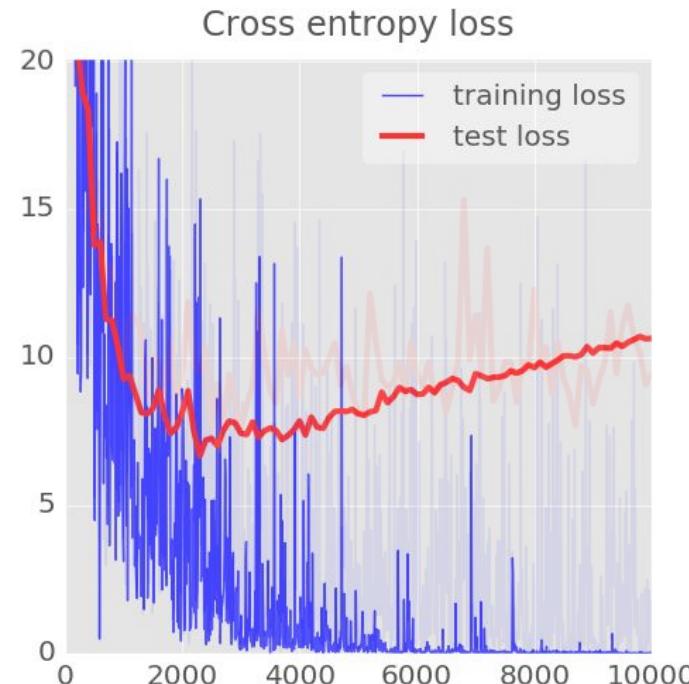
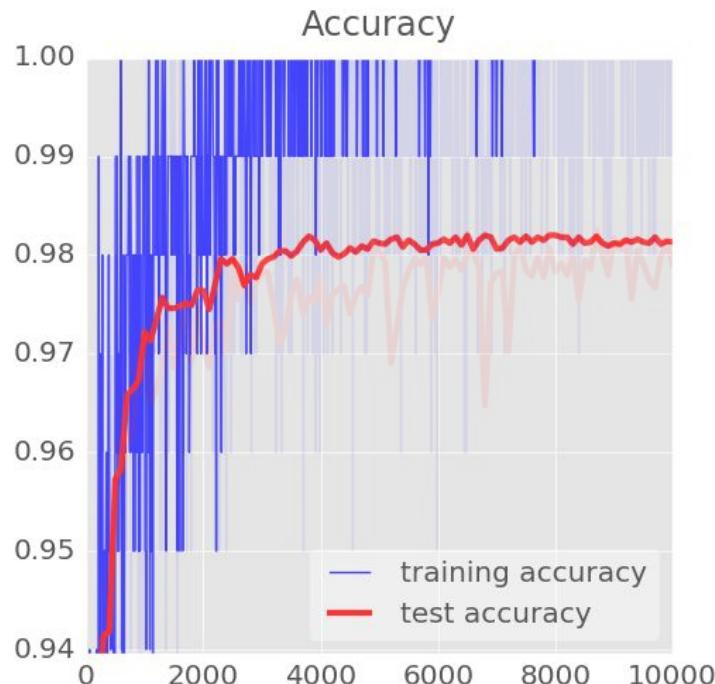


Demo



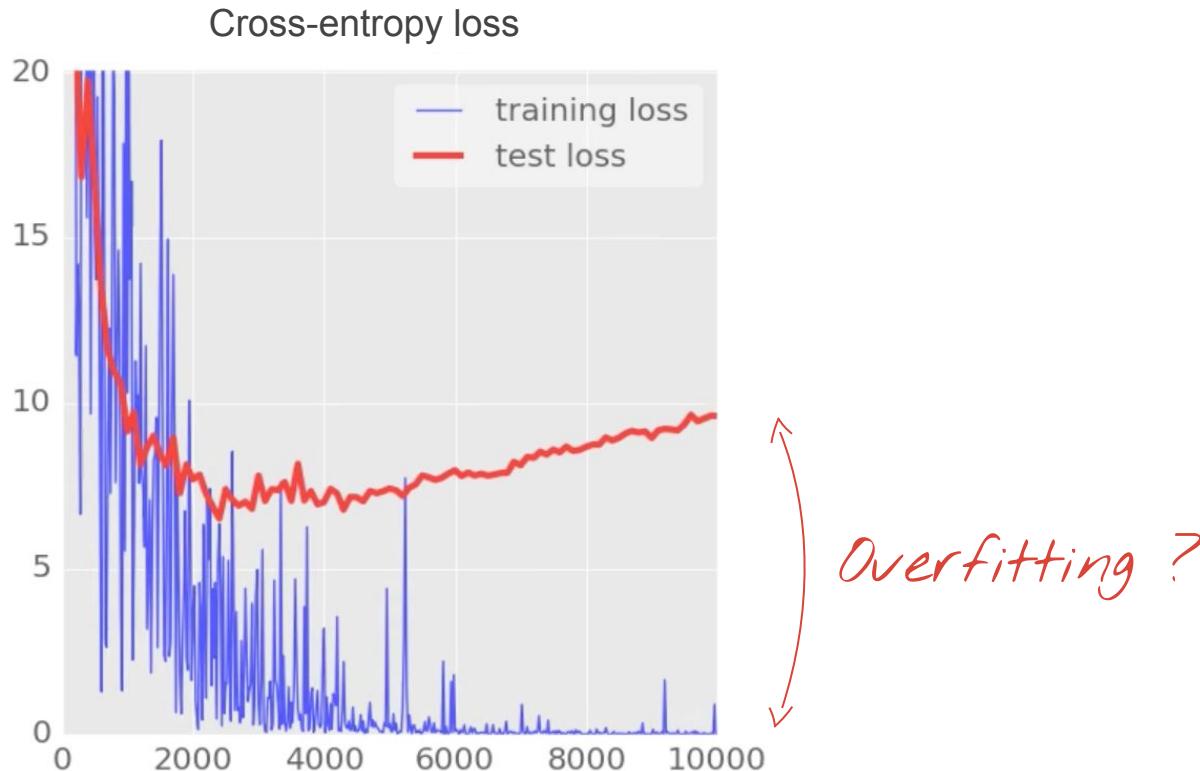


Learning rate decay



Learning rate 0.003 at start then dropping exponentially to 0.0001

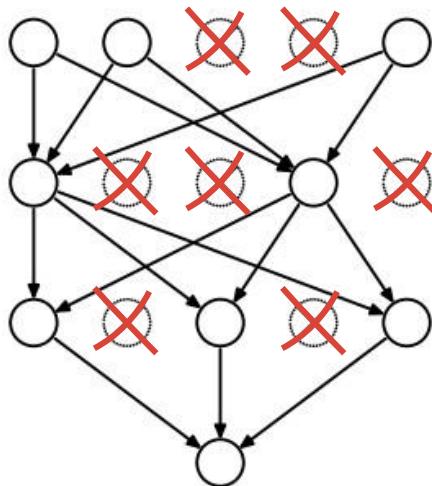
Overfitting



Dropout



Dropout

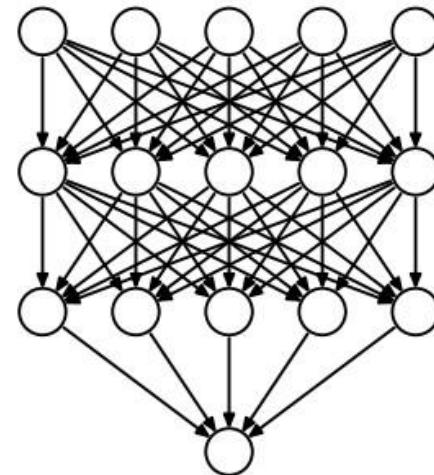


TRAINING
rate=0.5

```
pkeep =  
tf.placeholder(tf.float32)
```

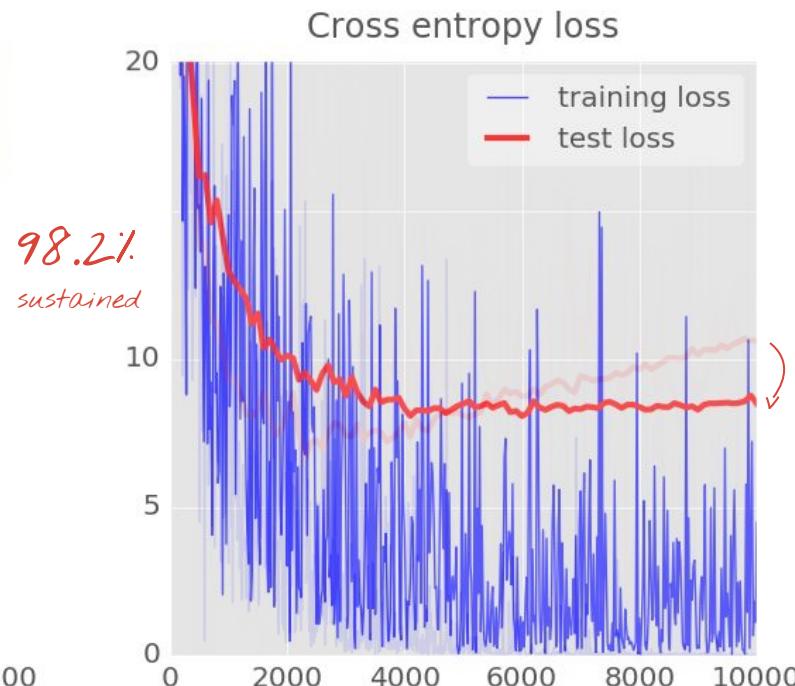
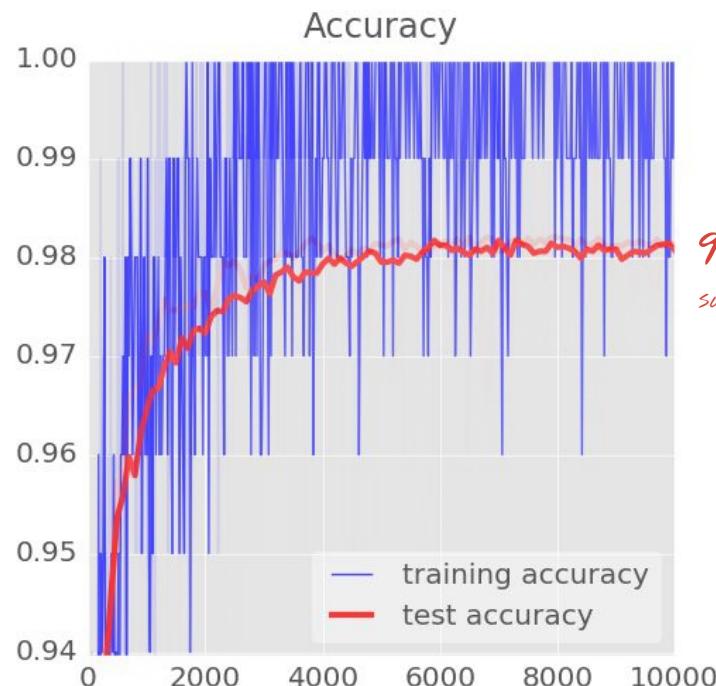
```
Yf = tf.nn.relu(tf.matmul(X, W) + B)
```

```
Y = tf.nn.dropout(Yf, pkeep)
```



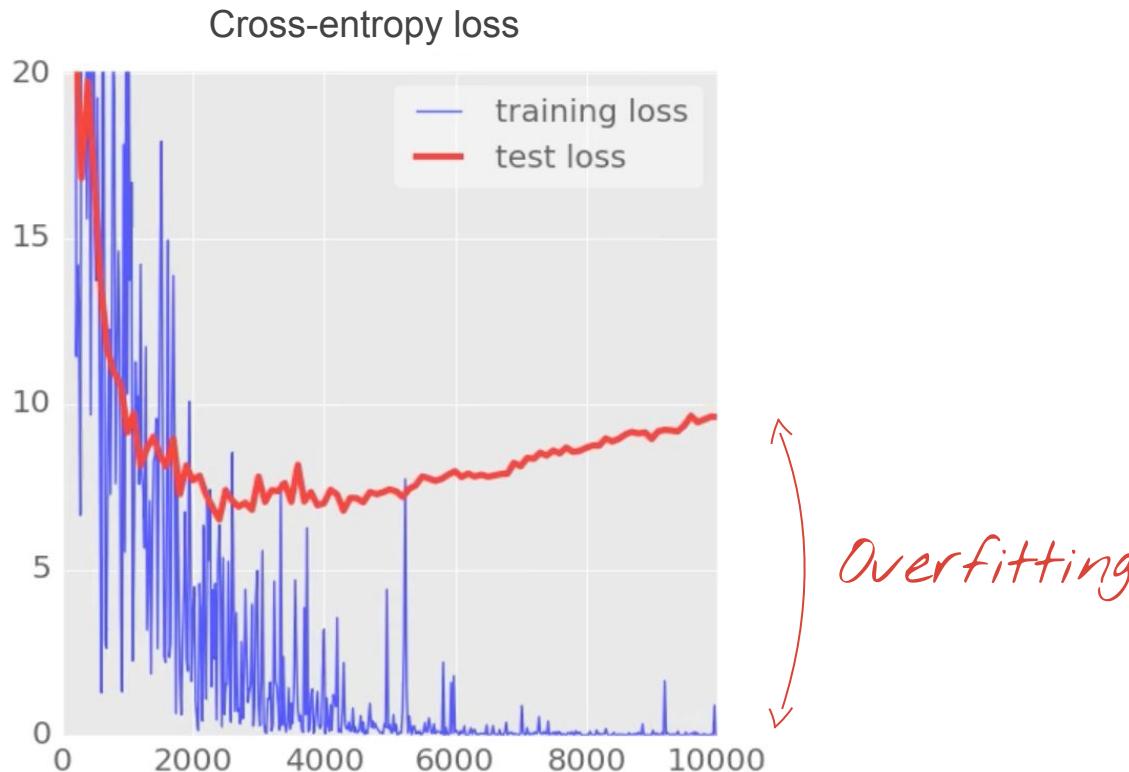
EVALUATION
rate=0

All the party tricks



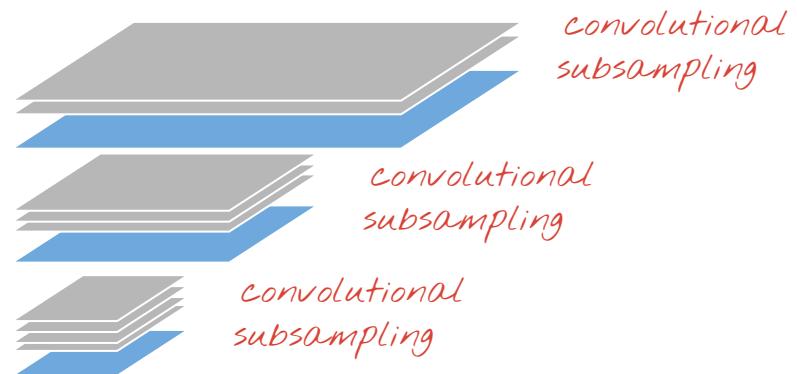
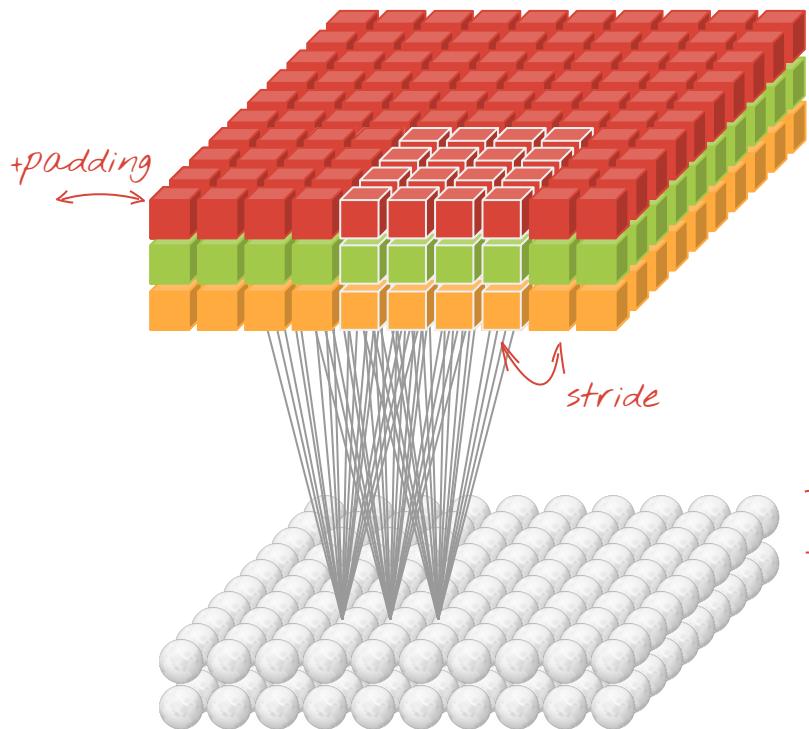
RELU, decaying learning rate $0.003 \rightarrow 0.0001$ and dropout 0.75

Overfitting





Convolutional layer



$W[4, 4, 3]$
 $W_2[4, 4, 3]$ | $W[4, 4, 3, 2]$

filter size input channels output channels

Hacker's tip

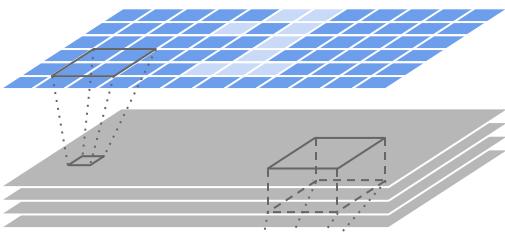


ALL
Convolu-
tional

Convolutional neural network

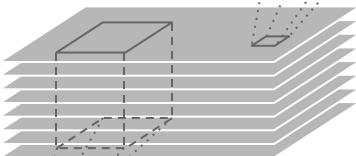
+ biases on
all layers

$28 \times 28 \times 1$

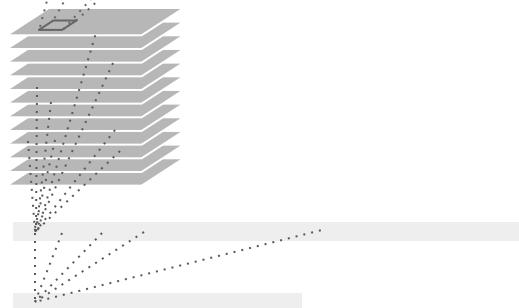


$28 \times 28 \times 4$

$14 \times 14 \times 8$



$7 \times 7 \times 12$



200

10

convolutional layer, 4 channels

$W1[5, 5, 1, 4]$ stride 1

convolutional layer, 8 channels

$W2[4, 4, 4, 8]$ stride 2

convolutional layer, 12 channels

$W3[4, 4, 8, 12]$ stride 2

fully connected layer

$W4[7 \times 7 \times 12, 200]$

softmax readout layer

$W5[200, 10]$

Tensorflow - initialisation

K=4

L=8

M=12

*filter
size* *input
channels* *output
channels*

W1 = tf.Variable(tf.truncated_normal([5, 5, 1, K], stddev=0.1))

B1 = tf.Variable(tf.ones([K])/10)

W2 = tf.Variable(tf.truncated_normal([5, 5, K, L], stddev=0.1))

B2 = tf.Variable(tf.ones([L])/10)

W3 = tf.Variable(tf.truncated_normal([4, 4, L, M], stddev=0.1))

B3 = tf.Variable(tf.ones([M])/10)

N=200

*weights initialised
with random values*

W4 = tf.Variable(tf.truncated_normal([7*7*M, N], stddev=0.1))

B4 = tf.Variable(tf.ones([N])/10)

W5 = tf.Variable(tf.truncated_normal([N, 10], stddev=0.1))

B5 = tf.Variable(tf.zeros([10])/10)

Tensorflow - the model

input image batch
 $X[100, 28, 28, 1]$

weights

stride

biases

```
Y1 = tf.nn.relu(tf.nn.conv2d(X, W1, strides=[1, 1, 1, 1], padding='SAME') + B1)
Y2 = tf.nn.relu(tf.nn.conv2d(Y1, W2, strides=[1, 2, 2, 1], padding='SAME') + B2)
Y3 = tf.nn.relu(tf.nn.conv2d(Y2, W3, strides=[1, 2, 2, 1], padding='SAME') + B3)
```

```
YY = tf.reshape(Y3, shape=[-1, 7 * 7 * M])
```

```
Y4 = tf.nn.relu(tf.matmul(YY, W4) + B4)
```

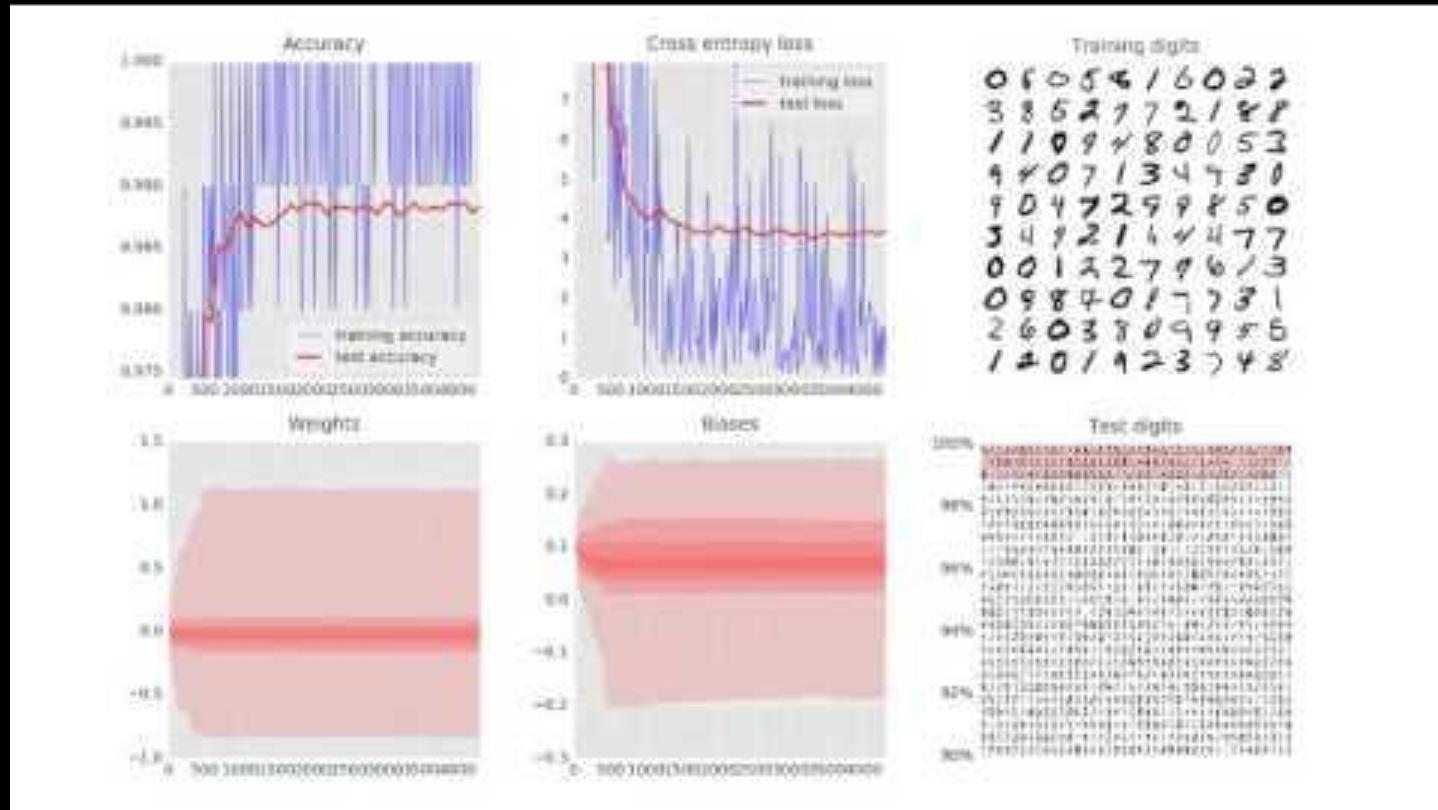
```
Y  = tf.nn.softmax(tf.matmul(Y4, W5) + B5)
```

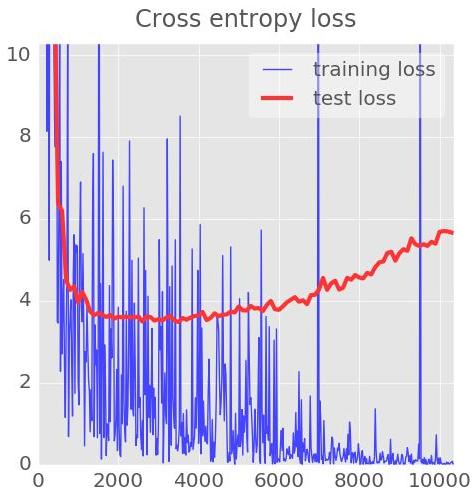
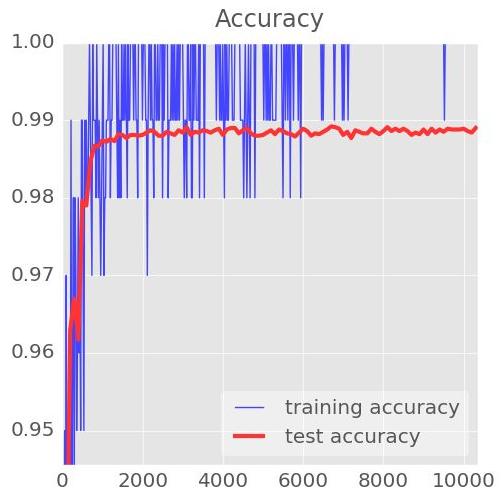
flatten all values for
fully connected layer

$Y3 [100, 7, 7, 12]$

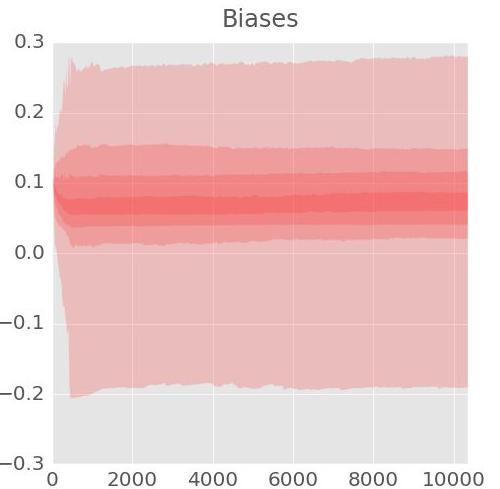
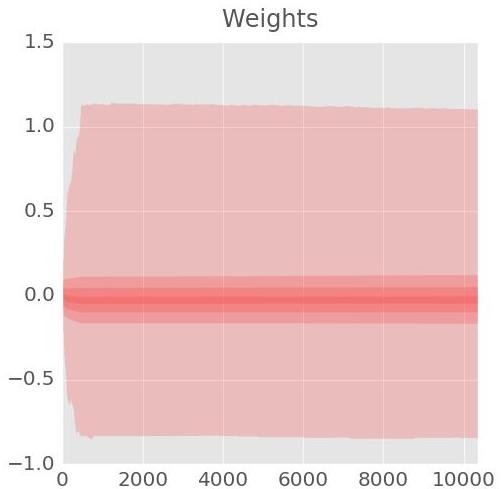
$YY [100, 7 \times 7 \times 12]$

Demo



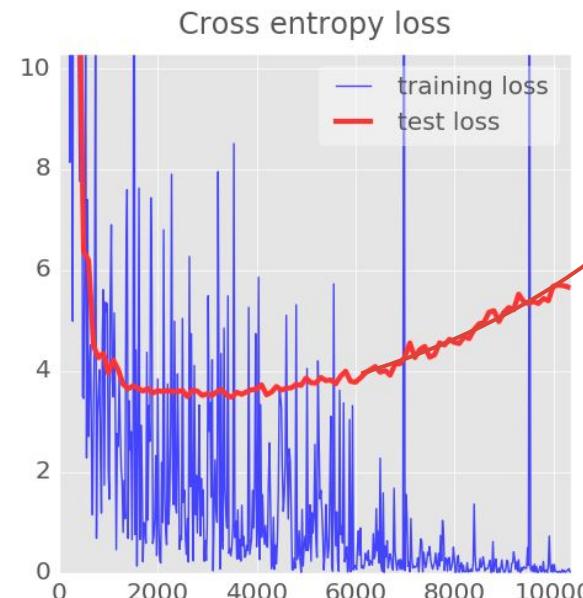
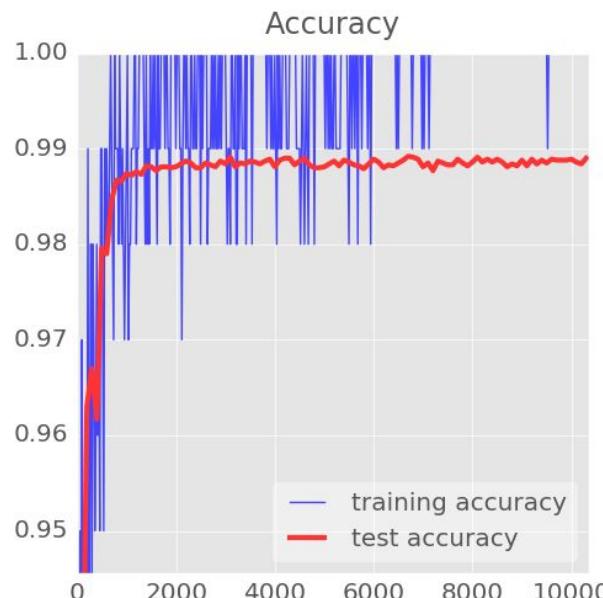


98.9%

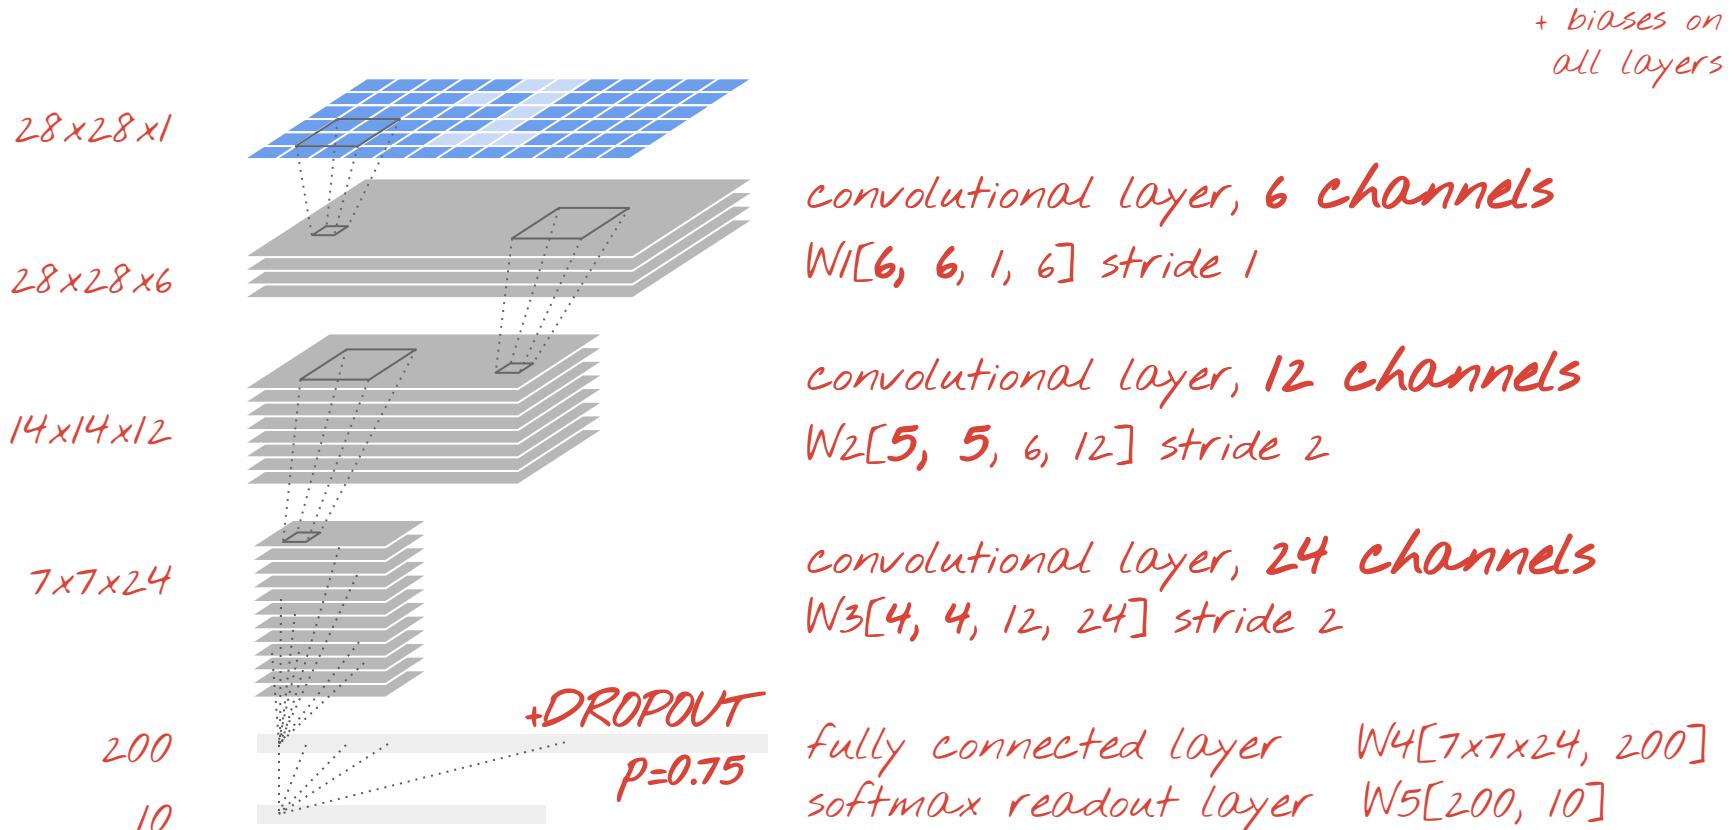


Test digits	
100%	42,3,4,1,5,6,1,4,1,1,1,7,0,3,0,3,0,9,4,7,3,1,8 15,1,0,5,1,1,2,9,0,-4,2,7,3,4,9,0,2,1,3,3,3,3,1,1,6,3,7,6,9,1 19,5,9,6,9,9,3,2,0,2,7,3,1,2,2,0,2,2,2,8,1,6,6,6,7,2,1,0,4,1,9,5,9 96,9,1,5,7,8,4,9,6,5,0,7,4,0,1,3,1,3,4,7,2,1,1,1,3,4,3,1,5,4 45,5,6,0,4,7,5,8,0,7,1,0,1,3,0,7,0,2,9,1,1,3,2,2,1,1,6,2,1,8,4,7,6,1 3,6,9,3,1,4,1,7,4,6,0,5,4,9,9,2,1,9,4,8,7,3,9,7,1,4,4,9,5,9,7,6,4,0,5,8,5 6,6,5,7,8,1,0,1,6,7,6,3,1,9,1,8,2,0,4,9,8,5,1,5,6,0,3,1,4,6,5,4,5,6,1,4 4,2,2,3,1,1,8,1,1,8,5,0,9,4,2,5,0,1,1,0,9,0,1,1,6,4,2,8,6,1,1,3,5,2,9 4,5,9,3,9,0,3,6,5,9,3,8,7,2,8,1,7,5,8,9,1,2,4,1,5,9,8,7,3,0,4,2,4 1,4,5,1,2,1,2,6,2,7,7,1,1,8,1,8,3,0,1,9,9,4,1,9,2,1,2,9,1,5,2,4,4,1,5,8 2,4,4,0,1,0,2,8,4,4,1,2,1,0,2,4,1,3,3,0,3,1,9,6,1,9,6,1,9,6,1,2,0,1,1,2 1,5,3,3,9,7,6,6,1,3,8,1,0,5,1,3,1,5,8,6,1,5,1,1,4,4,6,2,1,2,5,6,5,5,3,7,2,0 8,5,8,4,1,1,4,0,3,7,6,1,6,2,1,4,8,6,1,7,5,2,5,4,1,4,2,8,5,9,2,4,6,0,3,1,7,5 7,4,1,1,2,1,4,2,1,0,4,1,4,8,1,8,4,5,4,9,8,3,0,6,0,0,3,1,2,6,9,3,3,2,3 9,1,2,5,8,5,6,6,6,7,8,2,7,8,3,8,1,1,8,1,2,3,1,1,9,7,4,0,8,8,1,0,5,2,3 7,8,1,0,5,9,3,2,1,3,1,6,5,7,7,2,0,6,3,2,6,5,7,1,1,7,5,0,3,5,2,1,4,3,4 6,4,2,1,2,7,5,4,3,8,4,0,0,2,3,2,9,1,0,8,7,4,1,9,0,4,9,3,3,0,5,6,5,1,9,3 3,3,3,2,7,7,0,2,1,7,0,6,5,4,3,0,9,1,6,9,0,1,6,9,1,6,9,5,7,8,6,0,2,4,2,2 3,1,9,7,5,1,0,8,4,2,7,9,3,9,8,3,9,2,7,3,5,9,1,6,0,2,0,5,1,3,1,7,6,1,2,6 2,0,3,1,6,4,0,1,8,0,7,1,4,1,9,4,9,1,7,3,4,7,6,9,1,3,7,2,3,3,6,2,9,5,8,1,1 4,4,3,1,0,7,0,1,9,1,4,4,8,5,4,0,8,2,1,6,4,0,4,4,1,1,1,3,2,6,2,5,6,9,3,1,4,6 2,5,1,2,6,2,1,1,3,4,1,0,5,4,3,1,1,7,4,9,9,1,6,4,0,2,4,5,1,6,4,1,1,9,2,4,1 5,5,7,8,3,1,5,6,8,3,9,4,1,3,8,0,2,1,2,6,3,4,0,8,8,8,3,1,7,5,9,6,3,2,6 1,3,6,0,7,2,1,7,1,2,2,7,0,3,1,7,4,6,1,2,4,6,1,7,1,5,0,2,3,1,3,1,0,1,7,0,3,5 2,1,6,6,9,1,8,3,6,2,2,5,6,0,9,2,1,2,8,8,8,3,2,4,9,0,0,6,6,3,2,1,3,2,9,8,0,9,5
98%	
96%	
94%	
92%	
90%	

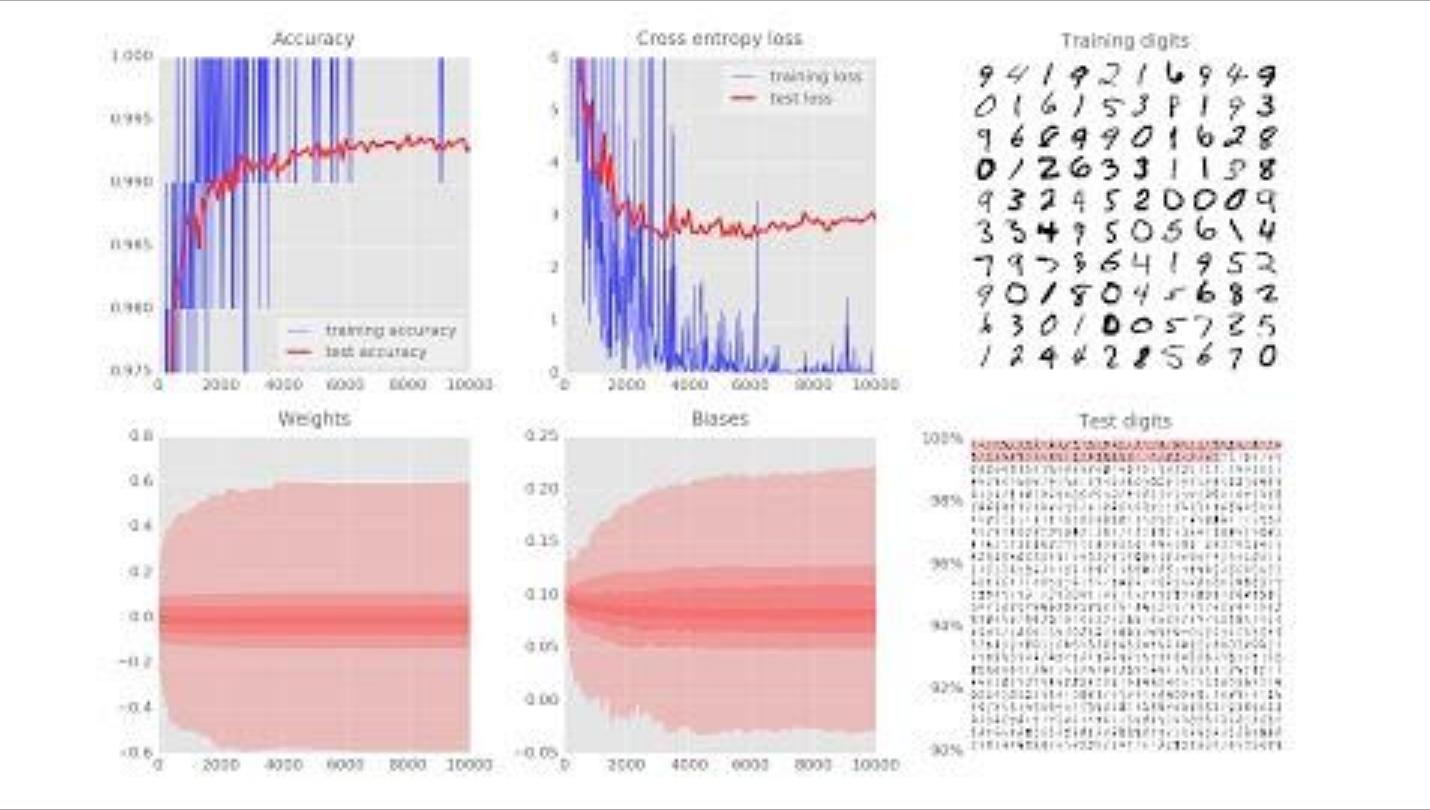
WTXH ???

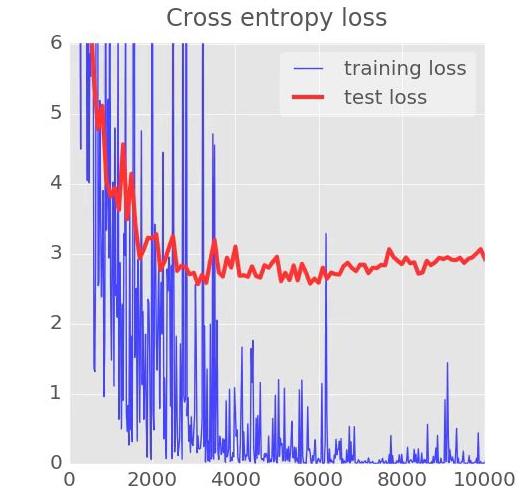
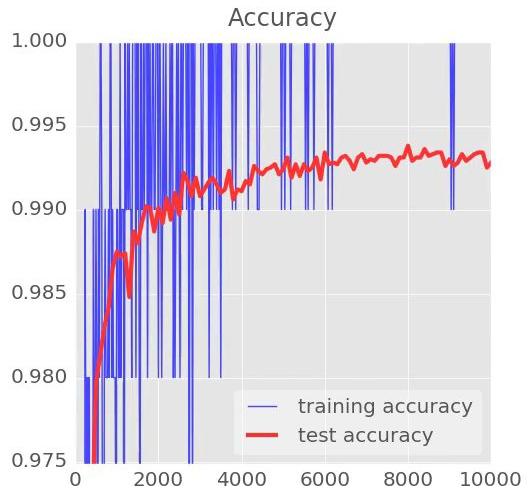


Bigger convolutional network + dropout

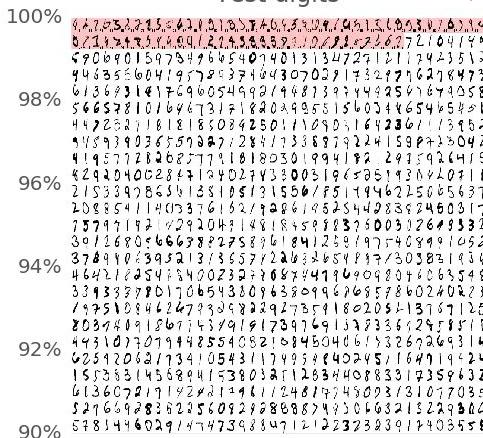
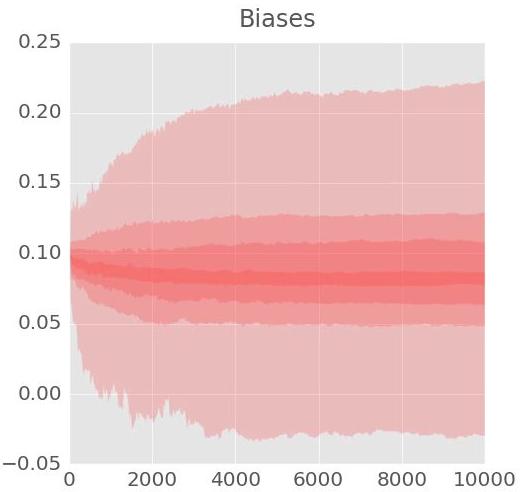
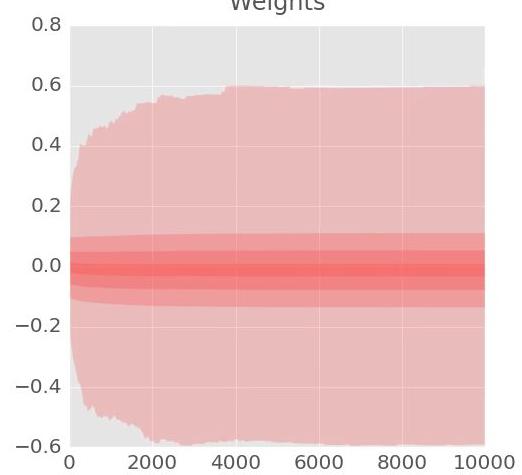


Demo

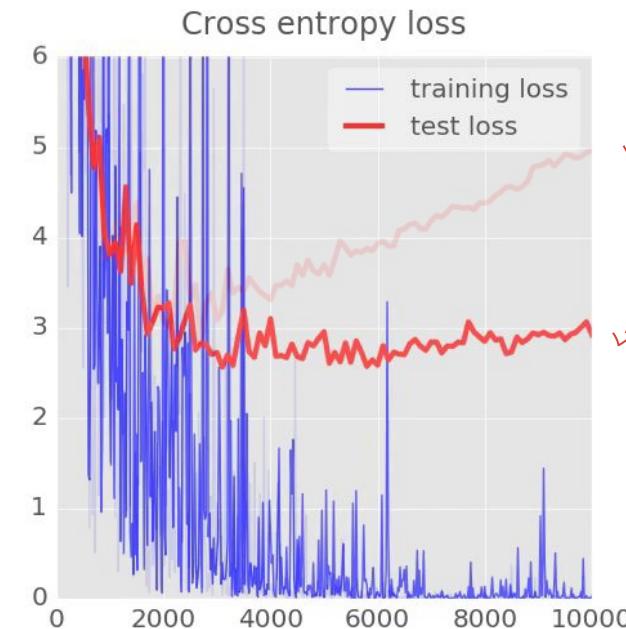
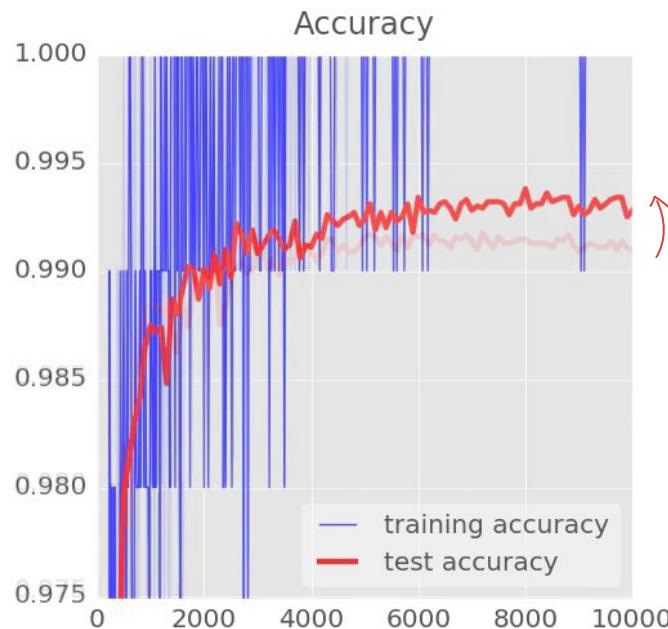




99.3%



YEAH !



with dropout



Slides Downloadable Link

https://docs.google.com/presentation/d/1B_ft3b1oLluOAOO2HzsvLwbN8fn1PzF4Kx44LF8P4Q0/edit?usp=sharing